

FREE PLAYBOOK

Gemini agent marketing prompt pack

Run a paid experiment end to end with one tight brief and seven minutes of review.

by Allen Anant Thomas, founder of The Growth Engine

READ TIME

~5 minutes. Read it once, apply it today.

WHY THIS MATTERS

The setup

Gemini 3 Pro can now run a full marketing test on its own. Pull the page, write variants, launch ads, watch significance, kill losers, write the memo. The catch is that a sloppy brief produces a fast wrong answer. The agent will optimize for whatever metric you name, even if that metric is the wrong one. This pack gives you the exact brief structure, the guardrails to attach, the metric hierarchy that prevents CTR traps, and the verification checklist to run before you accept the winner.

| 01 / 06

Pick one decision, not one experiment

Before opening Gemini, write the decision the test is supposed to settle. Not "test headlines" but "decide which of three buyer-intent angles, problem aware, solution aware, or competitor aware, books the most demos at or under \$80 CPA." The agent inherits your fuzziness. If you say "lift conversions" it picks the easiest signal it can measure inside the time window, which is usually CTR. Write the decision in one sentence with the variable, the population, the primary metric, and the cost ceiling. That sentence becomes the spine of every prompt below.

- Variable: the one thing that changes across variants.
- Population: who sees the test and on which page.
- Primary metric: the action that ties to revenue.
- Guardrail metric: what must not get worse.
- Cost ceiling: the dollar cap you accept losing.

| 02 / 06

Write the brief in Gemini agent mode

Open Gemini 3 Pro in agent mode and paste the brief below. The structure forces the agent to confirm scope, declare significance rules, and stop on guardrail breach instead of running until budget is gone. Always require it to print its plan before it spends a dollar. Read that plan. If the plan says it will optimize CTR alone, send it back. The 90 seconds you spend reading the plan is the difference between a trustworthy result and a 14-hour CTR mirage like the one in the post.

COPY-PASTE PROMPT

Role: senior growth marketer. Task: run a headline test on [URL]. Variable: H1 only. Variants: 6, mapped to problem aware, solution aware, and competitor aware intents, two each. Primary metric: demo bookings, tracked via [GA4 event or CRM webhook]. Guardrail: bounce rate must not rise more than 10 percent. Significance: 95 percent on demo rate, not CTR. Budget cap: \$200. Channel: Google Ads, exact match, brand excluded. Before launching, print the plan, the kill rules, and the UTM scheme. Wait for my approval.

| 03 / 06

Wire the right conversion event

The agent can only optimize what it can read. If your demo booking lives in HubSpot or Calendly and your ads platform only sees the button click, the agent will reward clicks and ignore booked meetings. Set up a server-side conversion in Google Ads using the Calendly or HubSpot webhook, or push a GA4 event named `demo_booked` from the thank-you page. Tell the agent the exact event name. Without this step, every prompt in this pack fails quietly. Verify the event fires by booking a fake demo yourself before the agent goes live.

- Google Ads: import GA4 conversion or use enhanced conversions.
- Calendly: native Google Ads integration sends booked events.
- HubSpot: use the Google Ads workflow action on form fill.

| 04 / 06

Force a hierarchy of metrics

Every brief needs three metrics ranked. Primary, guardrail, and tiebreaker. Primary is revenue adjacent, demo booked, trial started, paid. Guardrail is what must not regress, bounce rate, page time, lead quality score. Tiebreaker is used only when the primary cannot reach significance inside the budget, usually click through with a manual review flag. Tell the agent in plain text that the tiebreaker may not be used to declare a winner without a human signing off. This single instruction is what kills the CTR-only outcome that wasted Allen's first run.

COPY-PASTE PROMPT

Add to any agent brief: "Metric hierarchy. Primary: demo_booked, decision metric. Guardrail: bounce_rate, kill switch at +10 percent. Tiebreaker: CTR, advisory only, not a winner declaration. If primary metric cannot reach 95 percent significance within budget, return inconclusive plus a recommendation. Do not promote a tiebreaker winner without my written approval in this thread."

| 05 / 06

Verify the memo before you trust it

When the agent returns its 1-page memo, run a 7-minute audit. Open Google Ads and confirm spend matches the report. Open GA4 and check that the conversion count matches the memo. Pull the variant that won and read it out loud. If it sounds like a clickbait headline that overpromises, you probably bought a CTR win that will lose downstream. Check the bounce rate per variant. Check time on page. If the winner has the worst time on page, flag it and rerun with stricter guardrails before scaling.

- Spend reconciliation: Google Ads total vs memo total.
- Conversion match: GA4 demo_booked count vs memo count.
- Quality check: bounce rate and time on page per variant.
- Sanity read: would you click this and feel misled.

| 06 / 06

Reuse the template across channels

The same brief shape works for Meta Ads creative tests, email subject lines, and pricing page copy. Swap channel, variable, and event name. Keep the metric hierarchy and the kill rules identical. Save the master brief in a Notion or Google Doc and version it. Each time a test fails or surprises you, write one line under the brief explaining what broke and update the template. Within a month you will have a brief that produces clean results in any channel without rewriting from scratch.

COPY-PASTE PROMPT

Channel swap template: "Channel: [Meta Ads, email, on-site]. Variable: [creative, subject line, pricing copy]. Population: [audience or list segment]. Primary metric: [purchase, reply, plan_selected]. Guardrail: [unsubscribe rate, frequency, scroll depth]. Significance: 95 percent on primary. Budget or sample cap: [\$X or N sends]. Kill rules: stop any variant that crosses guardrail by 10 percent or burns 40 percent of budget with zero primary events."

RESOURCES

Tools and links you'll need

Gemini 3 Pro agent mode

Available inside the Gemini app for paid Advanced users. Toggle agent mode in the prompt bar.

Google Ads enhanced conversions

Setup guide at [support.google.com/google-ads, search enhanced conversions for leads](https://support.google.com/google-ads/search/enhanced%20conversions%20for%20leads).

GA4 conversion events

Create custom events at analytics.google.com under Admin, Events, Create event.

Calendly Google Ads integration

Native integration in Calendly settings, sends booked meetings as conversion events.

Statsig or GrowthBook

Free tier experimentation platforms that handle significance math when Google Ads cannot.

ONE LAST THING

The most common failure is trusting the memo because the agent sounds confident. Confidence is a writing style, not a statistical property. If the test ran for under 24 hours, treat any winner as a hypothesis, not a decision. Roll the winner to 50 percent of traffic for another 3 days before retiring the control. Speed is the gift. Verification is the rent.

Allen Anant Thomas

The Growth Engine

That's the playbook.

If this saved you time, the best thing you can do is share it. Tag a marketer who needs it.

WANT MORE LIKE THIS?

Connect on LinkedIn. I send a new free playbook every week.